



Module 09

Partha Pratim
Das

Objectives &
Outline

Additional Basic
Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

Database Management Systems

Module 09: Introduction to SQL/2

Partha Pratim Das

Department of Computer Science and Engineering
Indian Institute of Technology, Kharagpur

ppd@cse.iitkgp.ac.in



Module 09

Partha Pratim
Das

Objectives & Outline

Additional Basic Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

- Introduced relational query language
- Familiarized with data definition and basic query structure



Module 09

Partha Pratim
Das

Objectives & Outline

Additional Basic Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

- To complete the understanding of basic query structure



Module 09

Partha Pratim
Das

Objectives & Outline

Additional Basic Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

- Additional Basic Operations
 - Cartesian Product
 - Rename AS Operation
 - String Values
 - Order By
 - Select Top / Fetch
 - Where Clause Predicate
 - Duplicates



Module 09

Partha Pratim
Das

Objectives &
Outline

**Additional Basic
Operations**

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

Additional Basic Operations



Cartesian Product

Module 09

Partha Pratim
Das

Objectives &
Outline

Additional Basic
Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

- Find the Cartesian product *instructor X teaches*
select *
from *instructor, teaches*
 - generates every possible instructor-teaches pair, with all attributes from both relations
 - For common attributes (for example, *ID*), the attributes in the resulting table are renamed using the relation name (for example, *instructor.ID*)
- Cartesian product not very useful directly, but useful combined with where-clause condition (selection operation in relational algebra)



Cartesian Product

Module 09

Partha Pratim
DasObjectives &
OutlineAdditional Basic
Operations

Cartesian Product

Rename AS

Operation

String Values

Order By Clause

Select Top / Fetch
ClauseWhere Clause
Predicates

Duplicates

Module Summary

instructor

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821			
98345			

teaches

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009
32343	HIS-351	1	Spring	2010
45565	CS-101	1	Spring	2010
45565	CS-319	1	Spring	2010
76766	BIO-101	1	Summer	2009
				2010
				2009
				2009
				2010
				2010
				2009

Inst.ID	name	dept_name	salary	teaches.ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	12121	FIN-201	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	15151	MU-199	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	22222	PHY-101	1	Fall	2009
...
...
12121	Wu	Finance	90000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
12121	Wu	Finance	90000	10101	CS-347	1	Fall	2009
12121	Wu	Finance	90000	12121	FIN-201	1	Spring	2010
12121	Wu	Finance	90000	15151	MU-199	1	Spring	2010
12121	Wu	Finance	90000	22222	PHY-101	1	Fall	2009
...
...



Module 09

Partha Pratim
DasObjectives &
OutlineAdditional Basic
Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
ClauseWhere Clause
Predicates

Duplicates

Module Summary

- Find the names of all instructors who have taught some course and the course_id
 - select** *name, course_id*
 - from** *instructor, teaches*
 - where** *instructor.ID = teaches.ID*
 - Equi-Join, Natural Join

instructor				teaches					
ID	name	dept_name	salary	ID	course_id	sec_id	semester	year	
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009	
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010	
15151	Mozart	Music	40000	10101	CS-347	1	Fall	2009	
22222	Einstein	Physics	95000	12121	FIN-201	1	Spring	2010	
32343	El Said	History	60000	15151	MU-199	1	Spring	2010	
33456	Gold	Physics	87000	22222	PHY-101	1	Fall	2009	
45565	Katz	Comp. Sci.	75000	32343	HIS-351	1	Spring	2010	
58583	Califieri	History	62000	45565	CS-101	1	Spring	2010	
76543	Singh	Finance	80000	45565	CS-319	1	Spring	2010	
76766	Crick	Biology	72000	76766	BIO-101	1	Summer	2009	
83821								2010	
98345								2009	
inst.ID	name	dept_name	salary	teaches.ID	course_id	sec_id	semester	year	
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009	
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2010	
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2009	
10101	Srinivasan	Comp. Sci.	65000	12121	FIN-201	1	Spring	2010	
10101	Srinivasan	Comp. Sci.	65000	15151	MU-199	1	Spring	2010	
10101	Srinivasan	Comp. Sci.	65000	22222	PHY-101	1	Fall	2009	
...	
...	
12121	Wu	Finance	90000	10101	CS-101	1	Fall	2009	
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010	
12121	Wu	Finance	90000	10101	CS-347	1	Fall	2009	
12121	Wu	Finance	90000	12121	FIN-201	1	Spring	2010	
12121	Wu	Finance	90000	15151	MU-199	1	Spring	2010	
12121	Wu	Finance	90000	22222	PHY-101	1	Fall	2009	
...	
...	



Module 09

Partha Pratim
Das

Objectives &
Outline

Additional Basic
Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

- Find the names of all instructors in the Art department who have taught some course and the course_id

```
select name, course_id
from instructor, teaches
where instructor.ID = teaches.ID and instructor.dept_name = 'Art'
```



Rename AS Operation

Module 09

Partha Pratim
Das

Objectives &
Outline

Additional Basic
Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

- The SQL allows renaming relations and attributes using the **as** clause:
old_name as new_name
- Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.
select distinct *T.name*
from *instructor as T, instructor as S,*
where *T.salary > S.salary and S.dept_name = 'Comp. Sci'*
- Keyword **as** is optional and may be omitted
instructor as T \equiv *instructor T*



Cartesian Product Example

Module 09

Partha Pratim
Das

Objectives &
Outline

Additional Basic
Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

- Relation *emp_super*

<i>person</i>	<i>supervisor</i>
Bob	Alice
Mary	Susan
Alice	David
David	Mary

- Find the supervisor of “Bob”
- Find the supervisor of the supervisor of “Bob”
- Find ALL the supervisors (direct and indirect) of “Bob”



String Operations

Module 09

Partha Pratim
Das

Objectives &
Outline

Additional Basic
Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

- SQL includes a string-matching operator for comparisons on character strings. The operator **like** uses patterns that are described using two special characters:
 - percent (%). The % character matches any substring
 - underscore (_). The _ character matches any character
- Find the names of all instructors whose name includes the substring “dar”

```
select name  
from instructor  
where name like '%dar%'
```
- Match the string “100%”

```
like '100%' escape '\'
```
- in that above we use backslash (\) as the escape character



String Operations (2)

Module 09

Partha Pratim
Das

Objectives &
Outline

Additional Basic
Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

- Patterns are case sensitive
- Pattern matching examples:
 - 'Intro%' matches any string beginning with "Intro"
 - '%Comp%' matches any string containing "Comp" as a substring
 - ' _ _ ' matches any string of exactly three characters
 - ' _ _ _ %' matches any string of at least three characters
- SQL supports a variety of string operations such as
 - concatenation (using "||")
 - converting from upper to lower case (and vice versa)
 - finding string length, extracting substrings, etc.



Ordering the Display of Tuples

Module 09

Partha Pratim
Das

Objectives &
Outline

Additional Basic
Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

- List in alphabetic order the names of all instructors

```
select distinct name  
from instructor  
order by name
```

- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.
 - Example: **order by** *name desc*
- Can sort on multiple attributes
 - Example: **order by** *dept_name, name*



Selecting Number of Tuples in Output

Module 09

Partha Pratim
Das

Objectives &
Outline

Additional Basic
Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

- The **Select Top** clause is used to specify the number of records to return
- The **Select Top** clause is useful on large tables with thousands of records. Returning a large number of records can impact performance

```
select top 10 distinct name  
from instructor
```

- Not all database systems support the SELECT TOP clause.
 - SQL Server & MS Access support **select top**
 - MySQL supports the **limit** clause
 - Oracle uses **fetch first n rows only** and **rownum**

```
select distinct name  
from instructor  
order by name  
fetch first 10 rows only
```



Where Clause Predicates

Module 09

Partha Pratim
Das

Objectives &
Outline

Additional Basic
Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

- SQL includes a **between** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is, \geq \$90,000 and \leq \$100,000)

```
select name  
from instructor  
where salary between 90000 and 100000
```

- Tuple comparison

```
select name, course_id  
from instructor, teaches  
where (instructor.ID, dept_name) = (teaches.ID, 'Biology');
```




In Operator

Module 09

Partha Pratim
Das

Objectives &
Outline

Additional Basic
Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

- The **in** operator allows you to specify multiple values in a **where** clause
- The **in** operator is a shorthand for multiple **or** conditions
select *name*
from *instructor*
where *dept_name* **in** ('Comp. Sci.', 'Biology')



Duplicates

Module 09

Partha Pratim
Das

Objectives &
Outline

Additional Basic
Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

- In relations with duplicates, SQL can define how many copies of tuples appear in the result
- **Multiset** versions of some of the relational algebra operators – given multiset relations r_1 and r_2 :
 - a) $\sigma_\theta(r_1)$: If there are c_1 copies of tuple t_1 in r_1 , and t_1 satisfies selections σ_θ , then there are c_1 copies of t_1 in $\sigma_\theta(r_1)$
 - b) $\Pi_A(r)$: For each copy of tuple t_1 in r_1 , there is a copy of tuple $\Pi_A(t_1)$ in $\Pi_A(r_1)$ where $\Pi_A(t_1)$ denotes the projection of the single tuple t_1
 - c) $r_1 \times r_2$: If there are c_1 copies of tuple t_1 in r_1 and c_2 copies of tuple t_2 in r_2 , there are $c_1 \times c_2$ copies of the tuple $t_1.t_2$ in $r_1 \times r_2$



Duplicates (2)

Module 09

Partha Pratim
DasObjectives &
OutlineAdditional Basic
Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
ClauseWhere Clause
Predicates

Duplicates

Module Summary

- Example: Suppose multiset relations $r_1(A, B)$ and $r_2(C)$ are as follows:

$$r_1 = \{(1, a)(2, a)\} \quad r_2 = \{(2), (3), (3)\}$$

- Then $\Pi_B(r_1)$ would be $\{(a), (a)\}$, while $\Pi_B(r_1) \times r_2$ would be $\{(a, 2), (a, 2), (a, 3), (a, 3), (a, 3), (a, 3)\}$

- SQL duplicate semantics:

select A_1, A_2, \dots, A_n

from r_1, r_2, \dots, r_m

where P

is equivalent to the *multiset* version of the expression:

$$\Pi_{A_1, A_2, \dots, A_n}(\sigma_P(r_1 \times r_2 \times \dots \times r_m))$$



Module 09

Partha Pratim
Das

Objectives &
Outline

Additional Basic
Operations

Cartesian Product

Rename AS
Operation

String Values

Order By Clause

Select Top / Fetch
Clause

Where Clause
Predicates

Duplicates

Module Summary

- Completed the understanding of basic query structure

Slides used in this presentation are borrowed from <http://db-book.com/> with kind permission of the authors.

Edited and new slides are marked with “PPD”.