



Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

Database Management Systems

Module 49: Concurrency Control/1

Partha Pratim Das

Department of Computer Science and Engineering
Indian Institute of Technology, Kharagpur

ppd@cse.iitkgp.ac.in



Module 49

Partha Pratim
Das

Objectives & Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- With proper planning, a database can be recovered back to a consistent state from inconsistent state in the face of system failures. Such a recovery is done via cascaded or cascadeless rollback
- View Serializability is a weaker serializability system for better concurrency. However, testing for view serializability is NP complete



Module 49

Partha Pratim
Das

Objectives & Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- Concurrency Control through design of serializable schedule is difficult in general. Hence we take a look into locking mechanism and Lock-Based Protocols
- We need to understand how locks may be implemented



Module 49

Partha Pratim
Das

Objectives & Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- Concurrency Control
- Lock-Based Protocols
- Implementing Locking



Module 49

Partha Pratim
Das

Objectives &
Outline

**Concurrency
Control**

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

Concurrency Control



Concurrency Control

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- A database must provide a mechanism that will ensure that all possible schedules are both:
 - **Conflict serializable**
 - **Recoverable and, preferably, Cascadeless**
- A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency
- Concurrency-control schemes tradeoff between the amount of concurrency they allow and the amount of overhead that they incur
- Testing a schedule for serializability *after* it has executed is a little too late!
 - Tests for serializability help us understand why a concurrency control protocol is correct
- **Goal:** *To develop concurrency control protocols that will assure serializability*



Concurrency Control (2)

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- One way to ensure isolation is to require that data items be *accessed in a mutually exclusive manner*; that is, while one transaction is accessing a data item, no other transaction can modify that data item
 - Should a transaction hold a lock on the whole database
 - ▷ Would lead to strictly serial schedules – very poor performance
- The most common method used to implement locking requirement is to allow a transaction to access a data item only if it is currently holding a **lock** on that item



Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

**Lock-Based
Protocols**

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

Lock-Based Protocols



Lock-Based Protocols

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- *A lock is a mechanism to control concurrent access to a data item*
- Data items can be locked in two modes:
 - a) *exclusive (X)* mode:
 - Data item can be *both read as well as written*
 - **X-lock** is requested using **lock-X** instruction
 - b) *shared (S)* mode:
 - Data item can *only be read*
 - **S-lock** is requested using **lock-S** instruction
- A transaction can unlock a data item Q by the **unlock**(Q) Instruction
- Lock requests are made to the concurrency-control manager by the programmer
- *Transaction can proceed only after request is granted*



Lock-Based Protocols (2): Lock Compatibility Matrix

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- **Lock-Compatibility Matrix:** A lock compatibility matrix is used which states whether a data item can be locked by two transactions at the same time
- Full compatibility matrix

State of the lock	Lock request type →	
	Shared	Exclusive
Unlock	Yes	Yes
Shared	Yes	No
Exclusive	No	No

- Abbreviated compatibility matrix

State of the lock	Lock request type →	
	Shared	Exclusive
Shared	Yes	No
Exclusive	No	No



Lock-Based Protocols (3)

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- *Requesting for / Granting of a Lock*

- A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions

- *Sharing a Lock*

- Any number of transactions can hold shared locks on an item
- But if any transaction holds an exclusive lock on the item no other transaction may hold any lock on the item

- *Waiting for a Lock*

- If a lock cannot be granted, the requesting transaction is made to wait till all incompatible locks held by other transactions have been released

- *Holding a Lock*

- A transaction must hold a lock on a data item as long as it accesses that item

- *Unlocking / Releasing a Lock*

- Transaction T_i may unlock a data item that it had locked at some earlier point
- It is not necessarily desirable for a transaction to unlock a data item immediately after its final access of that data item, since serializability may not be ensured



Lock-Based Protocols: Example: Serial Schedule

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- Let A and B be two accounts that are accessed by transactions T_1 and T_2 .

- Transaction T_1 transfers \$50 from account B to account A
- Transaction T_2 displays the total amount of money in accounts A and B , that is, the sum $A + B$

- Suppose that the values of accounts A and B are \$100 and \$200, respectively
- If these transactions are executed serially, either as T_1, T_2 or the order T_2, T_1 then transaction T_2 will display the value \$300

T_1 :

```
lock-X(B);  
read(B);  
 $B := B - 50$ ;  
write(B);  
unlock(B);  
lock-X(A);  
read(A);  
 $A := A + 50$ ;  
write(A);  
unlock(A);
```

T_2 :

```
lock-S(A);  
read(A);  
unlock(A);  
lock-S(B);  
read(B);  
unlock(B);  
display( $A + B$ );
```



Lock-Based Protocols: Example (2): Concurrent Schedule: Bad

Module 49

Partha Pratim
DasObjectives &
OutlineConcurrency
ControlLock-Based
Protocols

Example

Lock-Based
ProtocolsTwo-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- If, however, these transactions are executed concurrently, then schedule 1 is possible
- In this case, transaction T_2 displays \$250, which is incorrect. The reason for this mistake is that
 - the transaction T_1 unlocked data item B too early, as a result of which T_2 saw an inconsistent state
- Suppose we delay unlocking till the end

T1:	T2:
lock-X(B);	lock-S(A);
read(B);	read(A);
$B := B - 50$;	unlock(A);
write(B);	lock-S(B);
unlock(B);	read(B);
lock-X(A);	unlock(B);
read(A);	display(A + B)
$A := A + 50$;	
write(A);	
unlock(A);	

T_1	T_2	concurrency control
lock-X(B)		grant-X(B, T_1)
read(B)		
$B := B - 50$		
write(B)		
unlock(B)	lock-S(A)	
	read(A)	grant-S(A, T_2)
	unlock(A)	
	lock-S(B)	
		grant-S(B, T_2)
	read(B)	
	unlock(B)	
	display(A + B)	
lock-X(A)		grant-X(A, T_1)
read(A)		
$A := A + 50$		
write(A)		
unlock(A)		

Schedule 1



Lock-Based Protocols: Example (3): Concurrent Schedule: Good

Module 49

Partha Pratim
DasObjectives &
OutlineConcurrency
ControlLock-Based
Protocols

Example

Lock-Based
ProtocolsTwo-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- Delaying unlocking till the end, T_1 becomes T_3 & T_2 becomes T_4

 T_3 :

```
lock-X(B);
read(B);
B := B - 50;
write(B);
lock-X(A);
read(A);
A := A + 50;
write(A);
unlock(B);
unlock(A)
```

 T_4 :

```
lock-S(A);
read(A);
lock-S(B);
read(B);
display(A + B);
unlock(A);
unlock(B)
```

- Hence, sequence of reads and writes as in Schedule 1 is no longer possible
- T_4 will correctly display \$300

T_3	T_4	concurrency control manager
lock-X(B)		grant-X(B, T_1)
read(B)		
$B := B - 50$		
write(B)		
unlock(B)		
	lock-S(A)	grant-S(A, T_2)
	read(A)	
	unlock(A)	
	lock-S(B)	grant-S(B, T_2)
	read(B)	
	unlock(B)	
	display(A + B)	
lock-X(A)		grant-X(A, T_1)
read(A)		
$A := A + 50$		
write(A)		
unlock(A)		

Schedule 1



Lock-Based Protocols: Example (4): Concurrent Schedule: Deadlock

Module 49

Partha Pratim Das

Objectives & Outline

Concurrency Control

Lock-Based Protocols

Example

Lock-Based Protocols

Two-Phase Locking Protocol

Lock Conversions

Automatic Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation of Locking

Lock Table

Module Summary

- Given, T_3 and T_4 , consider Schedule 2 (partial)
- Since T_3 is holding an exclusive mode lock on B and T_4 is requesting a shared-mode lock on B , T_4 is waiting for T_3 to unlock B
- Similarly, since T_4 is holding a shared-mode lock on A and T_3 is requesting an exclusive-mode lock on A , T_3 is waiting for T_4 to unlock A
- Thus, we have arrived at a state where neither of these transactions can ever proceed with its normal execution
- This situation is called **deadlock**
- When deadlock occurs, the system must roll back one of the two transactions.
- Once a transaction has been rolled back, the data items that were locked by that transaction are unlocked.
- These data items are then available to the other transaction, which can continue with its execution.

 T_3 :

```
lock-X(B);
read(B);
B := B - 50;
write(B);
lock-X(A);
read(A);
A := A + 50;
write(A);
unlock(B);
unlock(A)
```

 T_4 :

```
lock-S(A);
read(A);
lock-S(B);
read(B);
display(A + B);
unlock(A);
unlock(B)
```

T_3	T_4
lock-X(B) read(B) $B := B - 50$ write(B)	
lock-X(A)	lock-S(A) read(A) lock-S(B)

Schedule 2



Lock-Based Protocols

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- If we do not use locking, or if we unlock data items too soon after reading or writing them, we may get inconsistent states
- On the other hand, if we do not unlock a data item before requesting a lock on another data item, deadlocks may occur
- Deadlocks are a necessary evil associated with locking, if we want to avoid inconsistent states
- Deadlocks are definitely preferable to inconsistent states, since they can be handled by rolling back transactions, whereas inconsistent states may lead to real-world problems that cannot be handled by the database system
- A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks
- Locking protocols restrict the set of possible schedules
- The set of all such schedules is a proper subset of all possible serializable schedules
- We present locking protocols that allow only conflict-serializable schedules, and thereby ensure isolation



Two-Phase Locking Protocol

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- This protocol ensures conflict-serializable schedules
- Phase 1: Growing Phase
 - Transaction may obtain locks
 - Transaction may not release locks
- Phase 2: Shrinking Phase
 - Transaction may release locks
 - Transaction may not obtain locks
- The protocol assures serializability. It can be proved that the transactions can be serialized in the order of their **lock points**
 - That is, the point where a transaction acquired its final lock



Two-Phase Locking Protocol (2)

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

**Two-Phase Locking
Protocol**

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- There can be conflict serializable schedules that cannot be obtained if two-phase locking is used
- However, in the absence of extra information (that is, ordering of access to data), two-phase locking is needed for conflict serializability in the following sense:
 - Given a transaction T_i that does not follow two-phase locking, we can find a transaction T_j that uses two-phase locking, and a schedule for T_i and T_j that is not conflict serializable



Lock Conversions

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- Two-phase locking with lock conversions:
 - First Phase:
 - ▷ can acquire a lock- S on item
 - ▷ can acquire a lock- X on item
 - ▷ can convert a lock- S to a lock- X (upgrade)
 - Second Phase:
 - ▷ can release a lock- S
 - ▷ can release a lock- X
 - ▷ can convert a lock- X to a lock- S (downgrade)
- This protocol assures serializability. But still relies on the programmer to insert the various locking instructions



Automatic Acquisition of Locks: Read

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- A transaction T_i issues the standard read/write instruction, without explicit locking calls
- The operation **read**(D) is processed as:

if T_i has a lock on D

then

 read(D)

else begin

 if necessary, wait until no other transaction has a **lock-X** on D

 grant T_i a **lock-S** on D ;

 read(D)

end



Automatic Acquisition of Locks: Write

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- **write**(D) is processed as:
 - if** T_i has a **lock-X** on D
 - then**
 - write**(D)
 - else begin**
 - if necessary, wait until no other transaction has any lock on D ,
 - if T_i has a **lock-S** on D
 - then**
 - upgrade** lock on D to **lock-X**
 - else**
 - grant** T_i a **lock-X** on D
 - write**(D)
 - end;**
 - All locks are released after commit or abort



Deadlocks

Module 49

Partha Pratim
DasObjectives &
OutlineConcurrency
ControlLock-Based
Protocols

Example

Lock-Based
ProtocolsTwo-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- Two-phase locking *does not* ensure freedom from deadlocks

 T_3 :

```
lock-X(B);
read(B);
B := B - 50;
write(B);
lock-X(A);
read(A);
A := A + 50;
write(A);
unlock(B);
unlock(A)
```

 T_4 :

```
lock-S(A);
read(A);
lock-S(B);
read(B);
display(A + B);
unlock(A);
unlock(B)
```

T_3	T_4
lock-x (B)	lock-s (A) read (A) lock-s (B)
read (B)	
B := B - 50	
write (B)	
lock-x (A)	

- Observe that transactions T_3 and T_4 are two phase, but, in deadlock



Starvation

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- In addition to deadlocks, there is a possibility of **Starvation**
- **Starvation** occurs if the concurrency control manager is badly designed. For example:
 - A transaction may be waiting for an X-lock on an item, while a sequence of other transactions request and are granted an S-lock on the same item
 - The same transaction is repeatedly rolled back due to deadlocks
- Concurrency control manager can be designed to prevent starvation
- Starvation is also loosely referred to as **Livelock**



Cascading Rollback

Module 49

Partha Pratim
DasObjectives &
OutlineConcurrency
ControlLock-Based
Protocols

Example

Lock-Based
ProtocolsTwo-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- The potential for deadlock exists in most locking protocols. Deadlocks are a necessary evil
- When a deadlock occurs there is a possibility of cascading roll-backs
- Cascading roll-back is possible under two-phase locking
- In the schedule here, each transaction observes the two-phase locking protocol, but the failure of T5 after the read(A) step of T7 leads to cascading rollback of T6 and T7.

T_5	T_6	T_7
lock-X(A) read(A) lock-S(B) read(B) write(A) unlock(A)	lock-X(A) read(A) write(A) unlock(A)	lock-S(A) read(A)



More Two Phase Locking Protocols

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- To avoid Cascading roll-back, follow a modified protocol called **strict two-phase locking**
 - a transaction must hold all its exclusive locks till it commits/aborts
- **Rigorous two-phase locking** is even stricter
 - *All* locks are held till commit/abort. In this protocol transactions can be serialized in the order in which they commit
- Note that concurrency goes down as we move to more and more strict locking protocol



Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

Implementation of Locking



Implementation of Locking

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- A **lock manager** can be implemented as a separate process to which transactions send lock and unlock requests
- The lock manager replies to a lock request by sending a lock grant messages (or a message asking the transaction to roll back, in case of a deadlock)
- The requesting transaction waits until its request is answered
- The lock manager maintains a data-structure called a **lock table** to record granted locks and pending requests
- The lock table is usually implemented as an in-memory hash table indexed on the name of the data item being locked



Lock Table

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation


Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary



- Dark blue rectangles indicate granted locks; light blue indicate waiting requests
- Lock table also records the type of lock granted or requested
- New request is added to the end of the queue of requests for the data item, and granted if it is compatible with all earlier locks
- Unlock requests result in the request being deleted, and later requests are checked to see if they can now be granted
- If transaction aborts, all waiting or granted requests of the transaction are deleted
 - lock manager may keep a list of locks held by each transaction, to implement this efficiently



Module Summary

Module 49

Partha Pratim
Das

Objectives &
Outline

Concurrency
Control

Lock-Based
Protocols

Example

Lock-Based
Protocols

Two-Phase Locking
Protocol

Lock Conversions

Automatic
Acquisition of Locks

Deadlocks

Starvation

Cascading

More Protocols

Implementation
of Locking

Lock Table

Module Summary

- Understood the locking mechanism and protocols
- Realized that deadlock is a peril of locking and needs to be handled through rollback

Slides used in this presentation are borrowed from <http://db-book.com/> with kind permission of the authors.

Edited and new slides are marked with “PPD”.